

Program 3: Design and implement a Convolutional Neural Network (CNN) for classification of image dataset

AIM

To **design and implement a Convolutional Neural Network (CNN)** for the **classification of an image dataset**

Algorithm

Step 1: Import Libraries

Step 2: Load Dataset

Step 3: Preprocess the Data

Step 4: Build the CNN Model

Step 5: Compile the Model

Step 6: Train the Model

Step 7: Evaluate the Model

Step 8: Make Predictions

Step 9: Visualize Training Performance

Program:

```
# Import necessary libraries

import tensorflow as tf
from tensorflow.keras import models, layers

import matplotlib.pyplot as plt

import numpy as np

# MNIST dataset: 28x28 grayscale images of handwritten digits (0-9)
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

# Display the shape of dataset

print("Training data shape:", x_train.shape, y_train.shape)

print("Testing data shape:", x_test.shape, y_test.shape)
```

```
# Preprocess the data- Normalize pixel values from [0,255] to [0,1]
```

```
x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
# CNNs expect input with channels -> reshape to (samples, height, width, channels)
```

```
x_train = x_train.reshape(-1, 28, 28, 1)
```

```
x_test = x_test.reshape(-1, 28, 28, 1)
```

```
#Build CNN Model
```

```
model = models.Sequential([
```

```
    # Convolutional Layer 1
```

```
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
```

```
    layers.MaxPooling2D((2, 2)),
```

```
    # Convolutional Layer 2
```

```
    layers.Conv2D(64, (3, 3), activation='relu'),
```

```
    layers.MaxPooling2D((2, 2)),
```

```
    # Flatten the output from 2D to 1D
```

```
    layers.Flatten(),
```

```
    # Fully Connected Layer
```

```
    layers.Dense(64, activation='relu'),
```

```
    # Output Layer (10 classes for digits 0-9)
```

```
    layers.Dense(10, activation='softmax')
```

```
])
```

```
# Show model summary
```

```
model.summary()
```

```
# Step 4: Compile the Model
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
# Step 5: Train the Model
```

```
history = model.fit(x_train, y_train, epochs=3, validation_data=(x_test, y_test))
```

```
# Step 6: Evaluate Model Performance
```

```
loss, accuracy = model.evaluate(x_test, y_test, verbose=0)  
print(f"\nTest Accuracy: {accuracy:.4f}")
```

```
# Step 7: Make Predictions
```

```
predictions = model.predict(x_test[:5])
```

```
# Display first 5 test images, their true labels, and predicted labels
```

```
for i in range(5):  
    plt.imshow(x_test[i].reshape(28, 28), cmap="gray")  
    plt.title(f"True: {y_test[i]} | Predicted: {np.argmax(predictions[i])}")  
    plt.axis('off')  
    plt.show()
```

```
# Step 8: Plot Training History (Loss & Accuracy)
```

```
plt.figure(figsize=(12, 4))
```

```
# Plot training & validation loss
```

```
plt.subplot(1, 2, 1)

plt.plot(history.history['loss'], label='Training Loss')

plt.plot(history.history['val_loss'], label='Validation Loss')

plt.title('Loss Over Epochs')

plt.xlabel('Epoch')

plt.ylabel('Loss')

plt.legend()
```

```
# Plot training & validation accuracy

plt.subplot(1, 2, 2)

plt.plot(history.history['accuracy'], label='Training Accuracy')

plt.plot(history.history['val_accuracy'], label='Validation Accuracy')

plt.title('Accuracy Over Epochs')

plt.xlabel('Epoch')

plt.ylabel('Accuracy')

plt.legend()

plt.show()
```

```
Training data shape: (60000, 28, 28) (60000,)
Testing data shape: (10000, 28, 28) (10000,)
/usr/local/lib/python3.12/dist-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an `input_shape`/`input_dim`
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential_21"
```

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_3 (Conv2D)	(None, 11, 11, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_7 (Flatten)	(None, 1600)	0
dense_47 (Dense)	(None, 64)	102,464
dense_48 (Dense)	(None, 10)	650

Total params: 121,930 (476.29 KB)

Trainable params: 121,930 (476.29 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/3

1875/1875 ————— 32s 16ms/step - accuracy: 0.9033 - loss: 0.3094 - val_accuracy: 0.9859 - val_loss: 0.0438

Epoch 2/3

1875/1875 ————— 40s 16ms/step - accuracy: 0.9858 - loss: 0.0460 - val_accuracy: 0.9894 - val_loss: 0.0336

Epoch 3/3

1875/1875 ————— 41s 16ms/step - accuracy: 0.9895 - loss: 0.0330 - val_accuracy: 0.9896 - val_loss: 0.0291


Test Accuracy: 0.9896

1/1 ————— 0s 63ms/step


True: 2 | Predicted: 2




True: 1 | Predicted: 1



True: 7 | Predicted: 7



True: 0 | Predicted: 0



True: 4 | Predicted: 4

